

CS 5348 Operating Systems

Project 1: Shell Program

Version 1.0, 1st Sep 2023

Due date: 24th Sep 2023

This project is a group project. The group size is 2. It must be implemented in C. Sharing your work with others outside of your group is strictly prohibited. Do not copy code from any outside source. **Warning: if you plagiarize, you will be caught and reported.**

You can develop the project in csgrads1/cs3/csaws. Compile the code in any of the cs* machines. However, do not test/execute in the cs* machines. Use any of the {net01,..., net45} machines for program execution. You may create large number of processes due to a bug (more on how to handle forkbomb later). You do not want to slow down the cs* machines by executing in them. CSTech does not take kindly if you do this. **If you are caught executing your program in any of the cs* machines, you will be given a ticket. For every ticket you receive, we will deduct 5 points from your final score.**

Unix Shell

In this project, you will build a simple Unix shell. The shell is the heart of the command-line interface, and thus is central to the Unix/C programming environment. Mastering use of the shell is necessary to become proficient in this world; knowing how the shell itself is built is the focus of this project.

There are three specific objectives to this assignment:

- To further familiarize yourself with the Unix/Linux programming environment.
- To learn how processes are created, destroyed, and managed.
- To gain exposure to the necessary functionality in shells.

Overview

Your task is to implement a *command line interpreter (CLI)* or, as it is more commonly known, a *shell*. The shell should operate in this basic way: when you type in a command (in response to its prompt), the shell creates a child process that executes the command you entered and then prompts for more user input when it has finished.

The shells you implement will be like, but simpler than, the one you run in Unix. If you do not know what shell you are running, it's probably `bash`. One thing you should do on your own time is learn more about your shell, by reading the man pages or other online materials.

Program Specifications

Basic Shell: `dash`

Your basic shell, called `dash` (short for DALLas SHell), is basically an interactive loop: it repeatedly prints a prompt `dash>` (note the space after the greater-than sign), parses the input, executes the command specified on that line of input, and waits for the command to finish. This is repeated until the user types `exit`. The name of your final executable should be `dash`.

The shell can be invoked with either no arguments or a single argument; anything else is an error. Here is the no-argument way:

```
prompt> ./dash
dash>
```

At this point, `dash` is running, and ready to accept commands. Type away!

The mode above is called *interactive* mode. It allows the user to type commands directly. The shell also supports a *batch mode*, which instead reads input from a batch file and executes commands from therein. Here is how you run the shell with a batch file named `batch.txt`:

```
prompt> ./dash batch.txt
```

One difference between batch and interactive modes: in interactive mode, a prompt is printed (`dash>`). In batch mode, no prompt should be printed.

You should structure your shell such that it creates a process for each new command (the exception are *built-in commands*, discussed below). Your basic shell should be able to parse a command and run the program corresponding to the command. For example, if the user types `ls -la /tmp`, your shell should run the program `/bin/ls` with the given arguments `-la` and `/tmp` (how does the shell know that `ls` is in `/bin`? It's something called the shell **path**; more on this below).

Structure

Basic Shell

The shell is very simple (conceptually): it runs in a while loop, repeatedly asking for input to tell it what command to execute. It then executes that command. The loop continues indefinitely, until the user enters the built-in command `exit`, at which point it exits. That is it!

For reading lines of input, you should use `getline()`. This allows you to obtain arbitrarily long input lines with ease. Generally, the shell will be run in *interactive mode*, where the user types a command (one at a time) and the shell acts on it. However, your shell will also support *batch mode*, in which the shell is given an input file of commands; in this case, the shell should not read user input (from `stdin`) but rather from this file to get the commands to execute.

In either mode, if you hit the end-of-file marker (EOF), you should call `exit(0)` and exit gracefully.

To parse the input line into constituent pieces, you might want to use `strtok()` (or, if doing nested tokenization, use `strtok_r()`). Read the man page (carefully) for more details.

To execute commands, look into `fork()`, `exec()`, and `wait()/waitpid()`. See the man pages for these functions and read the relevant [book chapter](#) for a brief overview.

You will note that there are a variety of commands in the `exec` family; for this project, **you must use `execv`**, not any other call in the `exec` family. You should **not** use the `system()` library function call to run a command. Remember that if `execv()` is successful, it will not return; if it does return, there was an error (e.g., the command does not exist). The most challenging part is getting the arguments correctly specified.

Paths

In our example above, the user typed `ls` but the shell knew to execute the program `/bin/ls`. How does your shell know this?

It turns out that the user must specify a **path** variable to describe the set of directories to search for executables; the set of directories that comprise the path are sometimes called the *search path* of the shell. The path variable contains the list of all directories to search, in order, when the user types a command.

Important: Note that the shell itself does not *implement* `ls` or other commands (except built-ins). All it does is find those executables in one of the directories specified by path and create a new process to run them.

To check if a particular file exists in a directory and is executable, consider the `access()` system call. For example, when the user types `ls`, and `path` is set to include both `/bin` and `/usr/bin`, try `access("/bin/ls", X_OK)`. If that fails, try `"/usr/bin/ls"`. If that fails too, it is an error.

Your initial shell path should contain one directory: `/bin`

Note: Most shells allow you to specify a binary specifically without using a search path, using either **absolute paths** or **relative paths**. For example, a user could type the **absolute path** `/bin/ls` and execute the `ls` binary without a search path being needed. A user could also specify a **relative path** which starts with the current working directory and specifies the executable directly, e.g., `./main`. In this project, you **do not** have to worry about these features.

Built-in Commands

Whenever your shell accepts a command, it should check whether the command is a **built-in command** or not. If it is, it should not be executed like other programs. Instead, your shell will invoke your implementation of the built-in command. For example, to implement the `exit` built-in command, you simply call `exit(0)`; in your dash source code, which then will exit the shell.

In this project, you should implement `exit`, `cd`, and `path` as built-in commands.

- `exit`: When the user types `exit`, your shell should simply call the `exit` system call with 0 as a parameter. It is an error to pass any arguments to `exit`.
- `cd`: `cd` always takes one argument (0 or >1 args should be signaled as an error). To change directories, use the `chdir()` system call with the argument supplied by the user; if `chdir` fails, that is also an error.
- `path`: The `path` command takes zero or more arguments, with each argument separated by whitespace from the others. A typical usage would be like this: `dash> path /bin /usr/bin`, which would add `/bin` and `/usr/bin` to the search path of the shell. If the user sets `path` to be empty, then the shell should not be able to run any programs (except built-in commands). The `path` command always overwrites the old path with the newly specified path.

Redirection

Many times, a shell user prefers to send the output of a program to a file rather than to the screen. Usually, a shell provides this nice feature with the `>` character. Formally this is named as redirection of standard output. To make your shell users happy, your shell should also include this feature, but with a slight twist (explained below).

For example, if a user types, `ls -la /tmp > output`, nothing should be printed on the screen. Instead, the standard output of the `ls` program should be rerouted to the file `output`. In addition, the standard error output of the file should be rerouted to the file `output` (the twist is that this is a little different than standard redirection).

If the `output` file exists before you run your program, you should simply overwrite it (after truncating it).

The exact format of redirection is a command (and possibly some arguments) followed by the redirection symbol followed by a filename. Multiple redirection operators or multiple files to the right of the redirection sign are errors.

Note: don't worry about redirection for built-in commands (e.g., we will not test what happens when you type `path /bin > file`).

Parallel Commands

Your shell will also allow the user to launch parallel commands. This is accomplished with the ampersand operator as follows:

```
dash> cmd1 & cmd2 args1 args2 & cmd3 args1
```

In this case, instead of running `cmd1` and then waiting for it to finish, your shell should run `cmd1`, `cmd2`, and `cmd3` (each with whatever arguments the user has passed to it) in parallel, *before* waiting for any of them to complete.

Then, after starting all such processes, you must make sure to use `wait()` (or `waitpid`) to wait for them to complete. After all processes are done, return control to the user as usual (or, if in batch mode, move on to the next line).

Program Errors

You should print **this one and only error message** whenever you encounter an error of any type:

```
char error_message[30] = "An error has occurred\n";
write(STDERR_FILENO, error_message, strlen(error_message));
```

The error message should be printed to `stderr` (standard error), as shown above.

After most errors, your shell simply *continue processing* after printing the one and only error message. However, if the shell is invoked with more than one file, or if the shell is passed a bad batch file, it should exit by calling `exit(1)`.

There is a difference between errors that your shell catches and those that the program catches. Your shell should catch all the syntax errors specified in this project page. If the syntax of the command looks perfect, you simply run the specified program. If there are

any program-related errors (e.g., invalid arguments to `ls` when you run it, for example), the shell does not have to worry about that (rather, the program will print its own error messages and exit).

Miscellaneous Hints

Do not start coding immediately without thinking through the design. Write a simple high level design document first. Please bring it to me for a review if you wish. Do not be a trial-and-error programmer. Think, design, and code.

Remember to get the **basic functionality** of your shell working before worrying about all the error conditions and end cases. For example, first get a single command running (probably first a command with no arguments, such as `ls`).

Next, add built-in commands. Then, try working on redirection. Finally, think about parallel commands. Each of these requires a little more effort on parsing, but each should not be too hard to implement.

At some point, you should make sure your code is robust to white space of various kinds, including spaces () and tabs (`\t`). In general, the user should be able to put variable amounts of white space before and after commands, arguments, and various operators; however, the operators (redirection and parallel commands) do not require whitespace. Check the return codes of all system calls from the very beginning of your work. This will often catch errors in how you are invoking these new system calls. It's also just good programming practice.

Beat up your own code! You are the best (and in this case, the only) tester of this code. Throw lots of different inputs at it and make sure the shell behaves well. Good code comes through testing; you must run many different tests to make sure things work as desired. Don't be gentle -- other users certainly won't be.

Finally, keep versions of your code. More advanced programmers will use a source control system such as git. (If you are using git, make sure it is private.) Minimally, when you get a piece of functionality working, make a copy of your `.c` file (perhaps a subdirectory with a version number, such as `v1`, `v2`, etc.). By keeping older, working versions around, you can comfortably work on adding new functionality, safe in the knowledge you can always go back to an older, working version if need be.

Testing

Make sure you compile your program as follows:

```
gcc dash.c -o dash -Wall -Werror -O
```

There should not be any error messages and warning during compilation.

Sample testcases are in the directory /cs5348-xv6/src/testcases/P1 in csjaws.utdallas.edu. Look at README file first. Additional testcases may be provided later. You should also write your own test cases and test thoroughly. Writing good testcases is an important skill you need to do well in your career.

Your program may end up creating too many processes than intended due to a bug (fork bomb). To prevent yourself from being locked out of the system limit the number of processes you can create by adding this line to .bashrc file in your home directory.

```
ulimit -u 90
```

In case you reach the limit of 100, then login through a different terminal and kill dash and all the processes created by dash. Use commands "ps -u" and kill pid.

Hand-in Instructions

Name your program as dash.c. Copy your source code to the directory /cs5348-xv6/xxxxxyyyyy/P1 in csjaws, where xxxxyyyyy is your netid. (You will know more about csjaws in Exercise 1.)

If you have worked with a partner, only one of you need to submit the files. But both of you should create a text file named PARTNER in /cs5348-xv6/xxxxxyyyyy/P1 and save your partner's name and netid in the file.

Grading

Source code should be structured well with adequate comments clearly showing the different parts and functionalities implemented. 10% of the total points is allotted for style and comments, and the remaining 90% is for correctness.

The TA will compile the code and test it. *If your code does not compile, you will get 0 points. If you use any other system call other than `execv()` to execute a command, we will deduct 50 points from your final score.*

You will also be called by the TA to demonstrate your code. If you are not able to explain the working of your code, you will not be given any points.

If you have not contributed to this project, you will get 0 points.