Machine Learning Project 1

Scikit-Learn and Cross Validation Project

Lakshmi Chakradhar Vijayarao

Net ID: lxv230000

Introduction

In this project, I built classification models to check if the mushroom is edible or poisonous using the mushroom_mixed_50000.csv dataset. I used three machine learning algorithms: Decision Tree Classifier, Random Forest Classifier and KNN Classifier. The models were evaluated using grid search and cross-validation, and the final model was saved for accuracy competition.

Dataset

The dataset contains 50,000 samples with 20 features. The target variable has two classes: e (edible mushroom) and p (poisonous mushroom).

Model Building

I used the scikit-learn library to implement two models:

- Decision Tree Classifier
- Random Forest Classifier
- KNN Classifier

Cross-Validation and Grid Search

I performed grid search with 5-fold cross-validation. The parameter grids for each model were as follows:

Model	Parameter	Values
Decision Tree	max_depth	[10, 20, 30]
	min_samples_split	[2, 5, 10]
	min_samples_leaf	[1, 2, 4]
Random Forest	n_estimators	[50, 100]
	max_depth	[10, 15]
	min_samples_split	[5, 10]
	min_samples_leaf	[1, 2, 4]
	max_features	['sqrt', 'log2']
KNN	neighbors	[3. 5. 7]
	Weights.	['uniform' , 'distance']

Table 1: Grid search parameters for Decision Tree, Random Forest and KNN

Documentation

The approach for this project can be summarized in the following steps:

- **Data Loading**: I began by loading the dataset using pandas. The dataset was read into a DataFrame, where the features (X) were separated from the target variable (y).
- **Train-Test Split**: I split the data into training and testing sets using train test split from scikit-learn. This ensured that our model would be evaluated on unseen data, providing a fair assessment of its performance.
- **Model Definition**: I defined three classifiers: DecisionTreeClassifier, RandomForestClassifier and KNNClassifier.
- **Parameter Tuning**: To find the best hyperparameters, we employed Grid Search with cross-validation. This involved defining a parameter grid for each model, which allowed me to systematically explore different combinations of parameters to find the optimal settings for our classifiers.
- **Model Evaluation**: After fitting the models with the training data, we evaluated their performance using accuracy scores calculated on the test data. I also recorded the mean accuracy and standard deviation from the cross-validation results.
- **Model Saving**: The best performing model based on accuracy was saved for future use. I utilized Python's pickle module for serialization.

Results

Decision Tree:

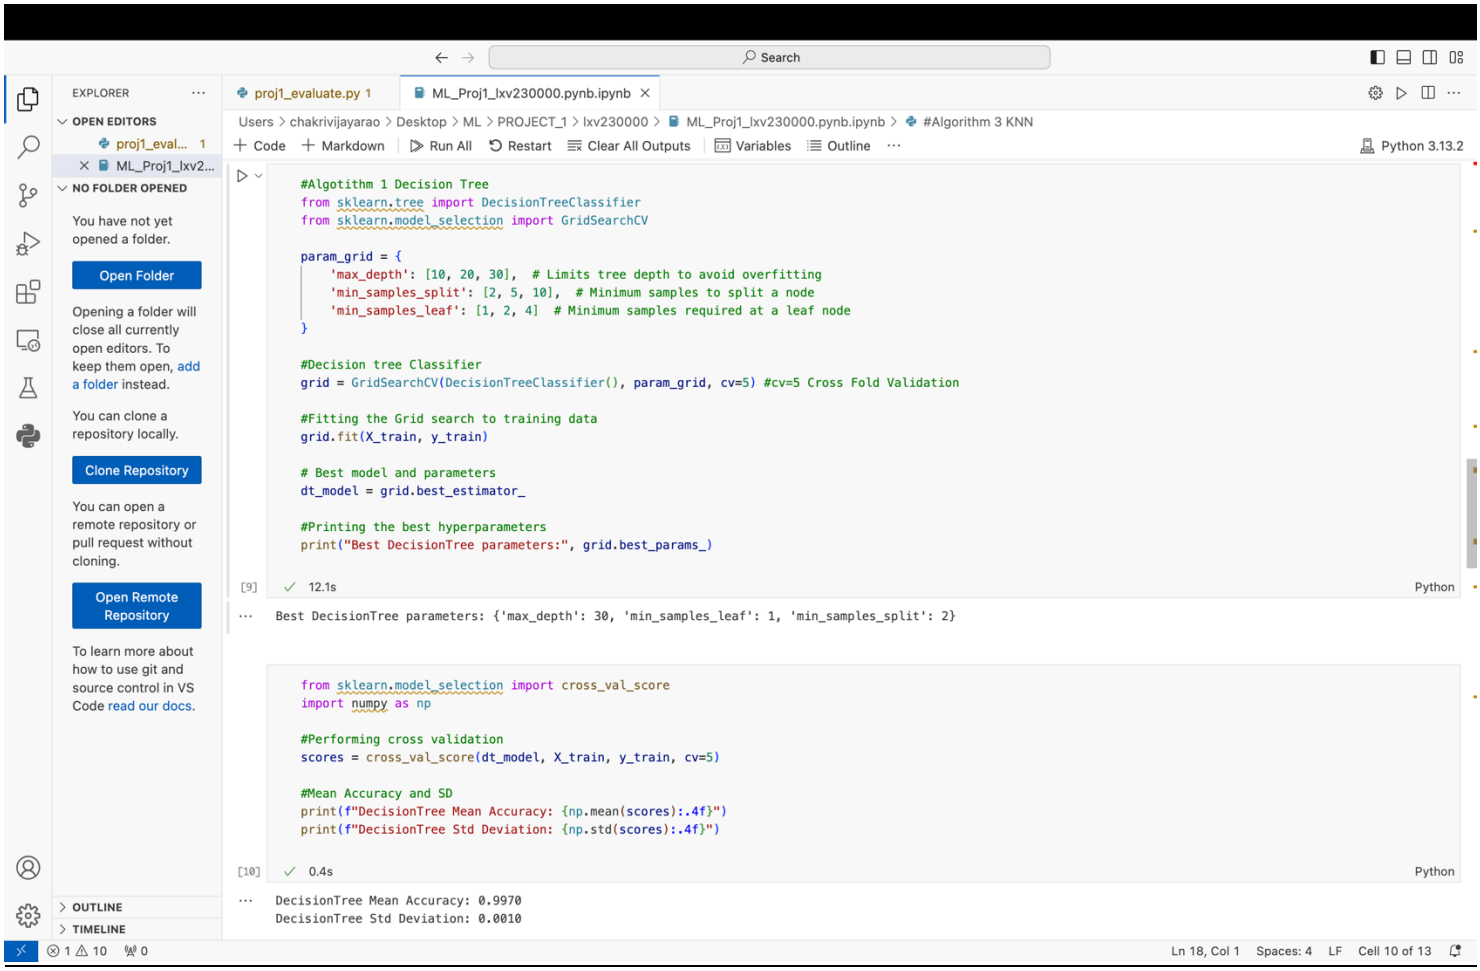

Random Forest:

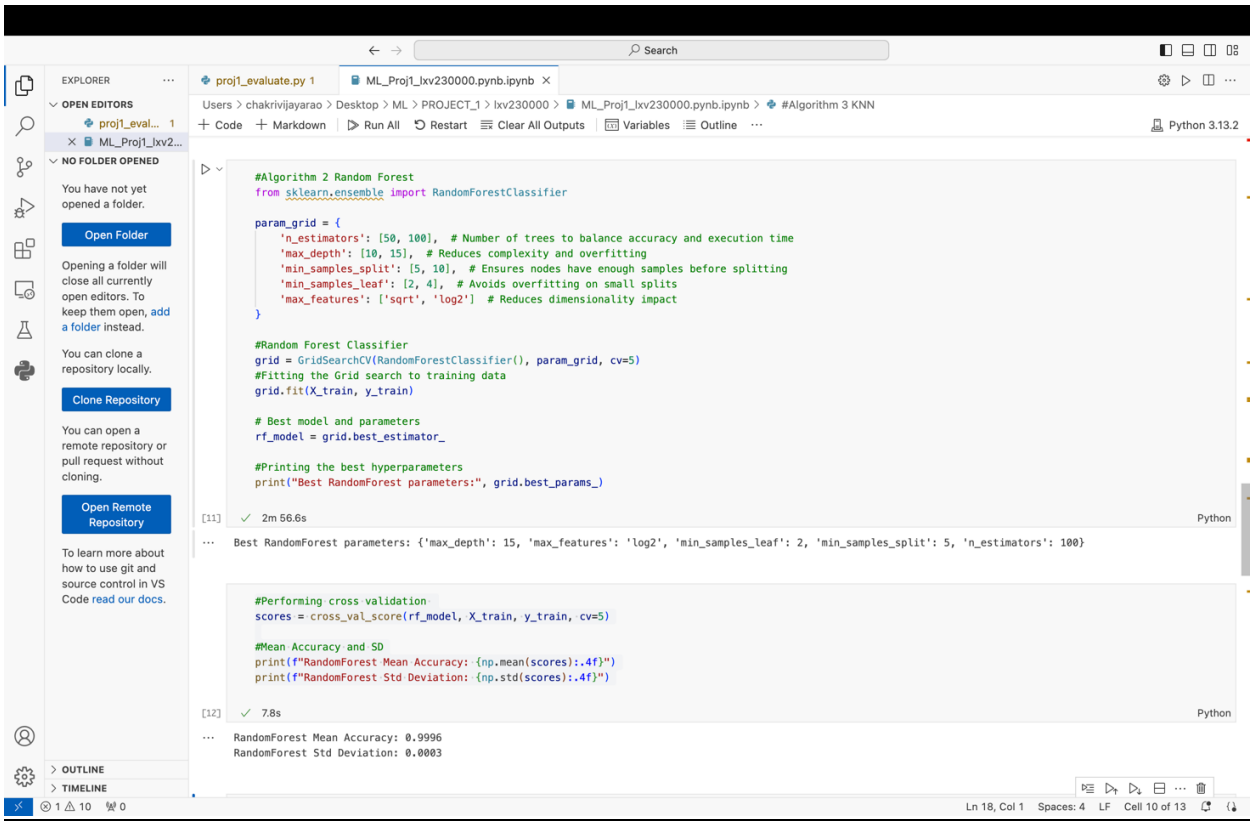

KNN:

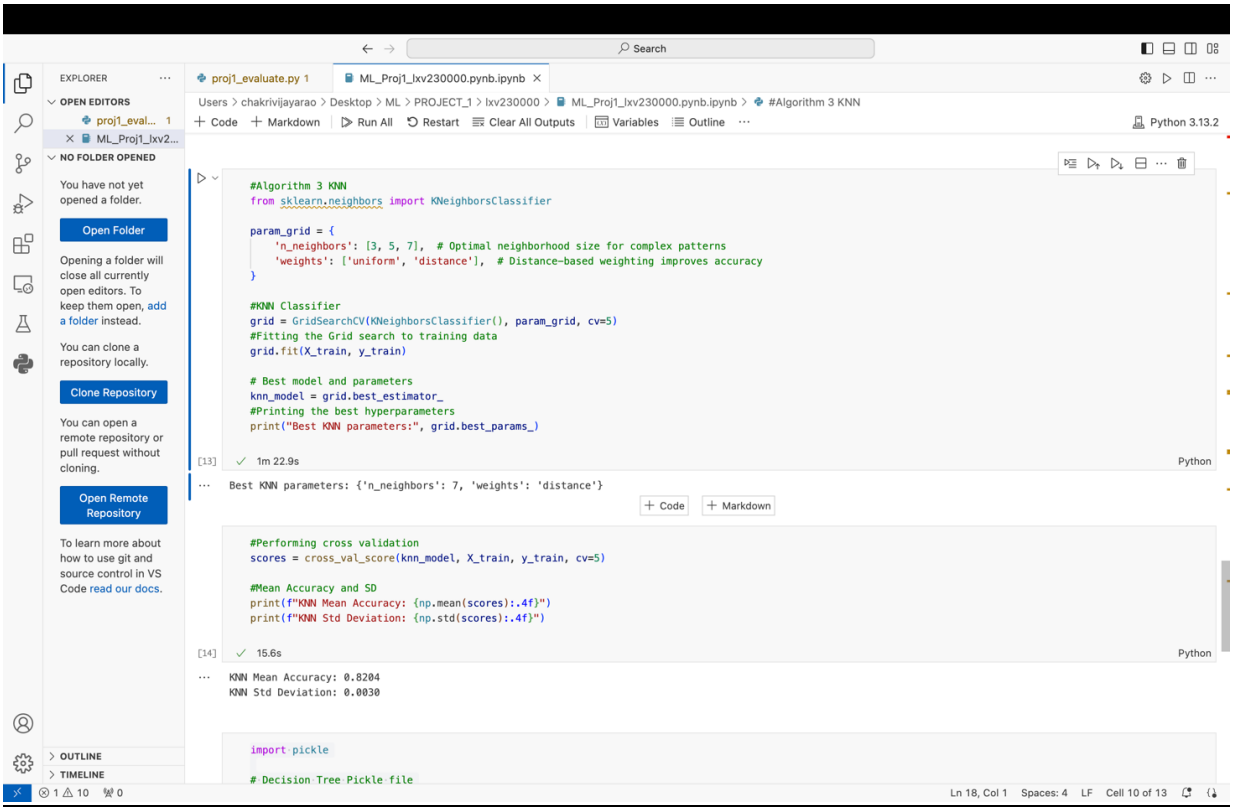

The results of grid search with cross-validation are summarized below:

Hyperparameter	Decision Tree Classifier	Random Forest Classifier	KNN Classifier
max depth	30	15	-
max features	-	Log2	-
min samples leaf	1	2	-
min samples split	2	5	-
n estimators	-	100	-
N neighbours	-	-	7
weights	-	-	distance

Table 2: Best Parameters Found for Decision Tree and Random Forest Classifiers

Model	Mean Accuracy	Accuracy Std. Dev.
Decision Tree	0.9970	0.0010
Random Forest	0.9996	0.0003
KNN	0.8204	0.0030

Table 3: Cross-Validation and Test Accuracy Results

Evalvuation Program Output:

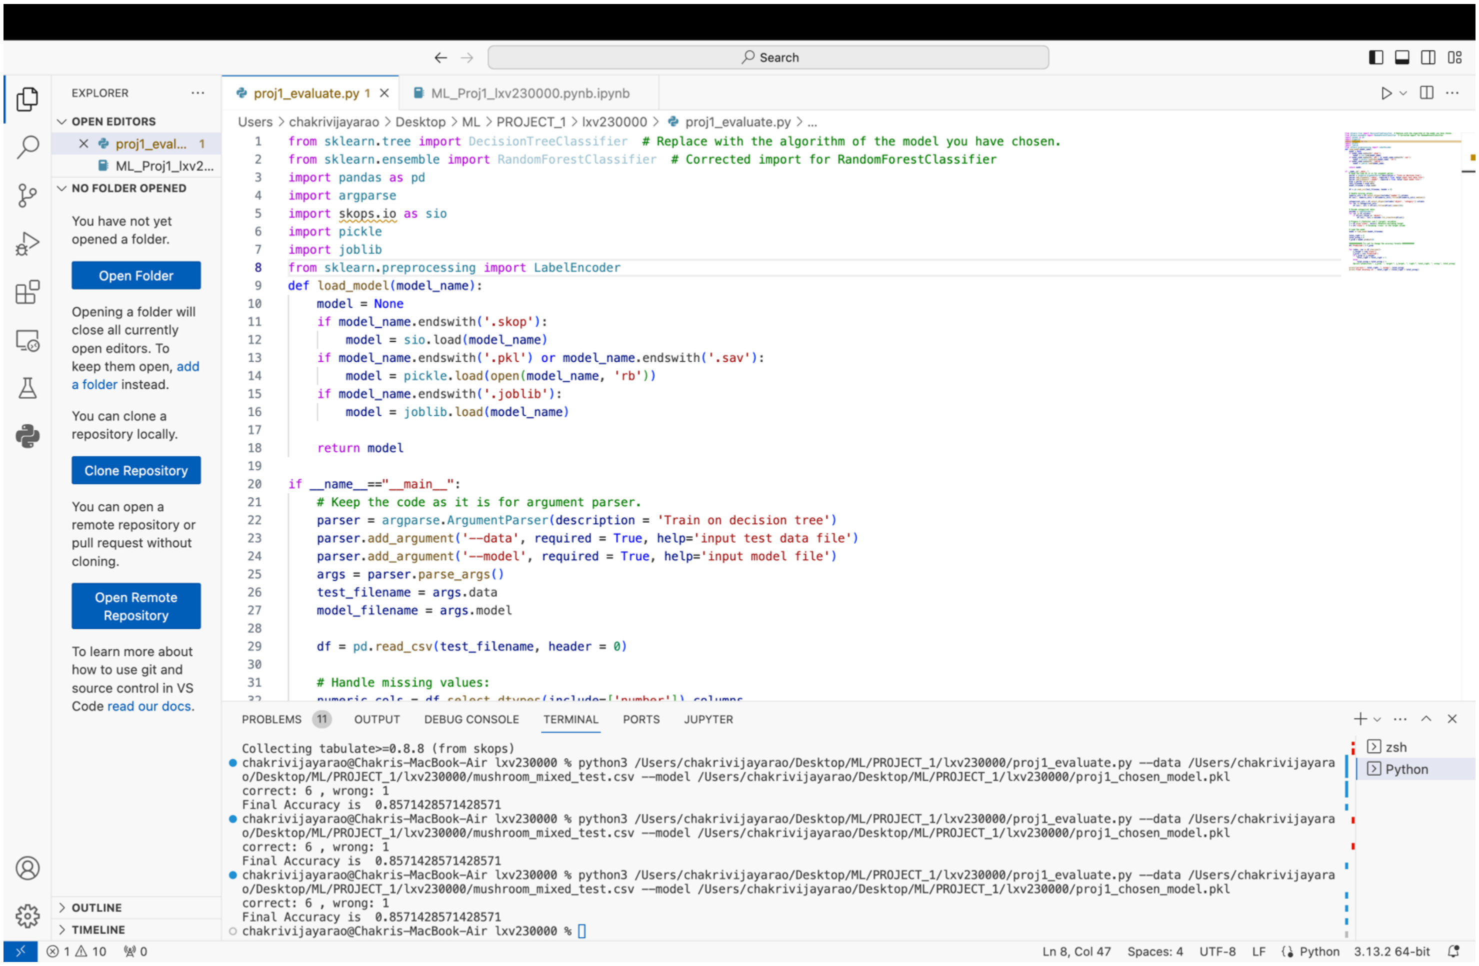

Conclusion

The Random Forest model achieved the highest accuracy of 99.96%, and it was saved as proj1_chosen_model.pkl.